

# Breast Cancer Prediction (Classification)

We have a dataset consisting of the information of 569 patients tested for breast cancer, we have 30 predictor variables based on the lumps caused by the tumor, and the response variable is the diagnosis: Benign or Malignant

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
sns.set()

df = pd.read_csv("breast-cancer.csv")

df.info()
df.head()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                569 non-null    float64
12  radius_se                             569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                               569 non-null    float64
16  smoothness_se                         569 non-null    float64
17  compactness_se                        569 non-null    float64
18  concavity_se                          569 non-null    float64
19  concave points_se                     569 non-null    float64
20  symmetry_se                           569 non-null    float64
21  fractal_dimension_se                  569 non-null    float64
22  radius_worst                          569 non-null    float64
23  texture_worst                         569 non-null    float64
24  perimeter_worst                       569 non-null    float64
25  area_worst                            569 non-null    float64
26  smoothness_worst                      569 non-null    float64
27  compactness_worst                     569 non-null    float64
28  concavity_worst                       569 non-null    float64
29  concave points_worst                   569 non-null    float64
30  symmetry_worst                        569 non-null    float64
31  fractal_dimension_worst                569 non-null    float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB
```

Out[1]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	...	radius_worst	texture_worst
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	...	25.38	17.33
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	...	24.99	23.41
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	...	23.57	25.53
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	...	14.91	26.50
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	...	22.54	16.67

5 rows × 32 columns

```
In [2]: fig = plt.figure(figsize=[10, 5])
sns.scatterplot(x="radius_mean", y="concave points_mean", hue="diagnosis", data=df)
plt.ylabel("mean for number of concave portions of the contour")
plt.xlabel("radius mean (mm)")
plt.legend(title="Diagnosis")
plt.show()
```



```
In [3]: def Bool(string):
    re = 0
    if string == "M":
        re = 1
    elif string != "B": #this shouldn't occur. We want to know if it happens, so we return nan
        re = np.nan #nan throws an error in regression models, so it will let us know if this happened
    return re
```

```
In [8]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression #random forest performed better than GBoost & logistic Regression
from sklearn.metrics import r2_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix

df_new = df.drop(columns = ['diagnosis'])
df_new['diagnosis'] = [Bool(i) for i in df['diagnosis']]

y = df_new["diagnosis"]
X = df_new.drop(columns = ['diagnosis', "id"]) #remove patient id

rf = RandomForestClassifier(random_state=42)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

rf.fit(X_train, y_train)
pred_tr = rf.predict(X_train)
pred_te = rf.predict(X_test)

r2_tr = r2_score(y_train, pred_tr)
r2_te = r2_score(y_test, pred_te)

print("r2 score for training set: ", r2_tr)
print("r2 score for testing set: ", r2_te, "\n")

print("accuracy score for training set: ", accuracy_score(y_train, pred_tr))
print("accuracy score for testing set: ", accuracy_score(y_test, pred_te), "\n")

print("confusion matrices are in the format: ")
print("[[benign,      predicted benign correctly      ",
      " benign,      predicted malignant falsely", " ]")
print(" [malignant, predicted benign falsely      ",
      " malignant, predicted malignant correctly]]", "\n")

print("training set confusion matrix")
print(confusion_matrix(y_train, pred_tr))

print("testing set confusion matrix")
print(confusion_matrix(y_test, pred_te))

r2 score for training set:  1.0
r2 score for testing set:  0.8512276321265085

accuracy score for training set:  1.0
accuracy score for testing set:  0.965034965034965

confusion matrices are in the format:
[[benign,      predicted benign correctly      benign,      predicted malignant falsely ]
 [malignant, predicted benign falsely      malignant, predicted malignant correctly]]

training set confusion matrix
[[268  0]
 [ 0 158]]
testing set confusion matrix
[[87  2]
 [ 3 51]]

The model results are good. on the testing data set, the random forest model predicted 96.5% correctly, and performing well in predicting both benign and malignant tumors, as seen in the confusion matrix.
```

With some more data (we only had 569 patients total), the model would likely perform even better.

```
In [ ]:
```